

# midterm-1

February 25, 2025

```
[37]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
df = pd.read_csv('student_enrollment_data.csv')
print(df.info())
print(df.describe())
print(df.columns)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2575 entries, 0 to 2574
Data columns (total 23 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Reference ID           2575 non-null   int64
1   Sex                    2575 non-null   object
2   Student Type           2575 non-null   object
3   Major 1                2561 non-null   object
4   Major 2                2163 non-null   object
5   IsPastorKid            2575 non-null   object
6   IsAthlete              2575 non-null   object
7   Race                   2280 non-null   object
8   Zip Code               2238 non-null   float64
9   AddressState           2497 non-null   object
10  Email_clicks_Jan        2575 non-null   int64
11  Email_click_Feb         2575 non-null   int64
12  Email_click_Mar         2575 non-null   int64
13  Email_click_Apr         2575 non-null   int64
14  Email_click_May         2575 non-null   int64
15  Email_click_Jun         2575 non-null   int64
16  Email_click_Jul         2575 non-null   int64
17  Email_click_Aug         2575 non-null   int64
18  Email_click_Sep         2575 non-null   int64
19  Email_click_Oct         2575 non-null   int64
20  Email_click_Nov         2575 non-null   int64
21  Email_click_Dec         2575 non-null   int64
22  Enrolled                2575 non-null   object
dtypes: float64(1), int64(13), object(9)
```

memory usage: 462.8+ KB

None

	Reference ID	Zip Code	Email_clicks_Jan	Email_click_Feb	\
count	2.575000e+03	2238.000000	2575.000000	2575.000000	
mean	4.770369e+08	21705.662198	1.126602	1.126214	
std	2.997064e+08	19710.784678	1.682758	1.651914	
min	1.551550e+05	745.000000	0.000000	0.000000	
25%	2.035259e+08	13750.750000	0.000000	0.000000	
50%	4.717699e+08	14623.000000	0.000000	0.000000	
75%	7.434801e+08	17884.250000	2.000000	2.000000	
max	9.997468e+08	99577.000000	12.000000	14.000000	

	Email_click_Mar	Email_click_Apr	Email_click_May	Email_click_Jun	\
count	2575.000000	2575.000000	2575.000000	2575.000000	
mean	1.075340	1.131262	0.766214	1.483495	
std	1.767947	1.954179	1.445630	2.240151	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	0.000000	
75%	2.000000	2.000000	1.000000	2.500000	
max	15.000000	16.000000	13.000000	17.000000	

	Email_click_Jul	Email_click_Aug	Email_click_Sep	Email_click_Oct	\
count	2575.000000	2575.000000	2575.000000	2575.000000	
mean	1.293592	1.099806	0.286214	0.622913	
std	1.864280	1.551417	0.709591	1.231049	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	0.000000	
75%	2.000000	2.000000	0.000000	1.000000	
max	13.000000	11.000000	10.000000	15.000000	

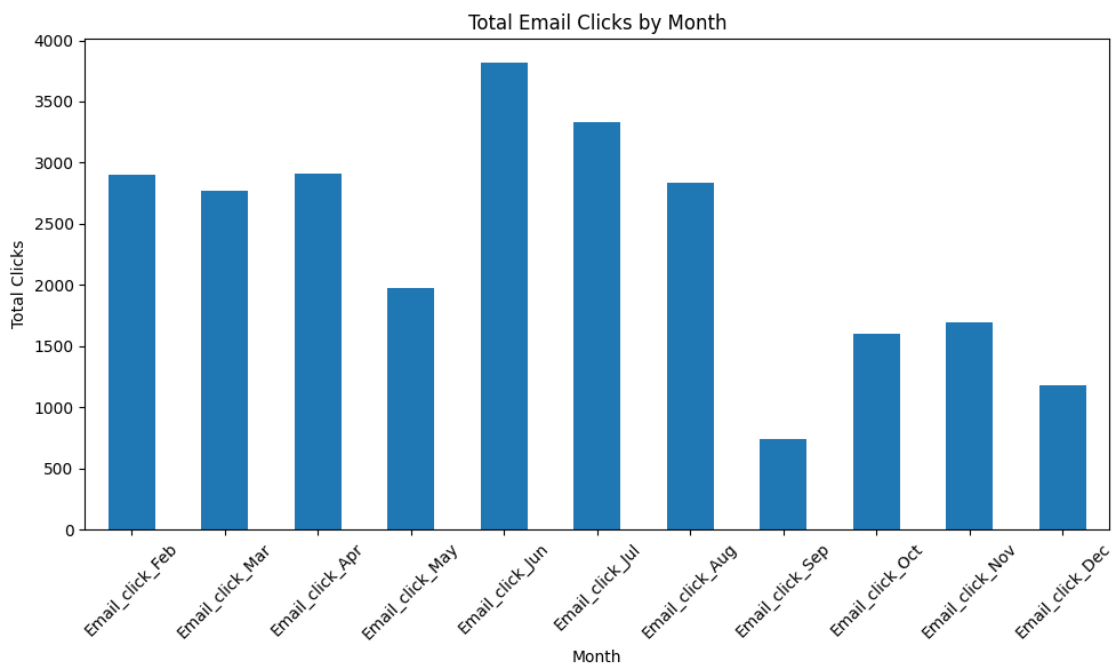
	Email_click_Nov	Email_click_Dec
count	2575.000000	2575.000000
mean	0.659417	0.458641
std	1.197366	0.977872
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	1.000000	1.000000
max	9.000000	11.000000

```
Index(['Reference ID', 'Sex', 'Student Type', 'Major 1', 'Major 2',
      'IsPastorKid', 'IsAthlete', 'Race', 'Zip Code', 'AddressState',
      'Email_clicks_Jan', 'Email_click_Feb', 'Email_click_Mar',
      'Email_click_Apr', 'Email_click_May', 'Email_click_Jun',
      'Email_click_Jul', 'Email_click_Aug', 'Email_click_Sep',
      'Email_click_Oct', 'Email_click_Nov', 'Email_click_Dec', 'Enrolled'],
      dtype='object')
```

```
[38]: email_columns = [col for col in df.columns if col.startswith('Email_click_')]
total_clicks = df[email_columns].sum()

plt.figure(figsize=(10,6))
total_clicks.plot(kind='bar')
plt.title('Total Email Clicks by Month')
plt.xlabel('Month')
plt.ylabel('Total Clicks')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

print(f"Month with highest total clicks: {total_clicks.idxmax()}")
```

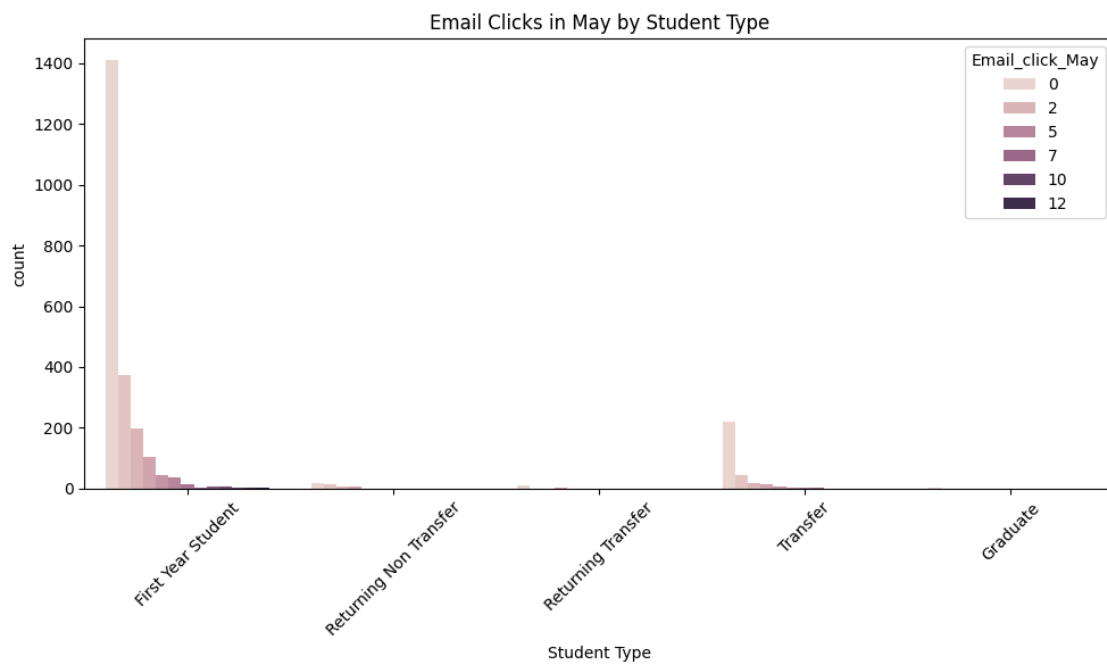
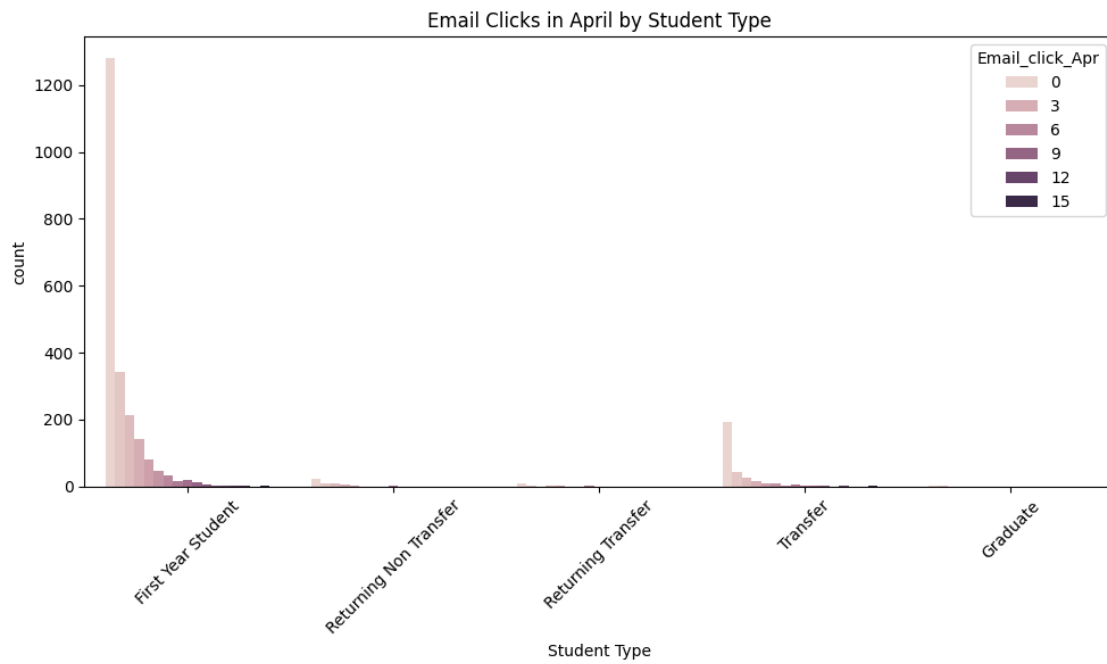


Month with highest total clicks: Email\_click\_Jun

```
[39]: plt.figure(figsize=(10,6))
sns.countplot(x='Student Type', hue='Email_click_Apr', data=df)
plt.title('Email Clicks in April by Student Type')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

plt.figure(figsize=(10,6))
sns.countplot(x='Student Type', hue='Email_click_May', data=df)
```

```
plt.title('Email Clicks in May by Student Type')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



Find percentage of students who clicked on emails

```
[40]: def calculate_email_click_percentage(df, column):
        total_students = df.shape[0]
        email_clicks = df[column].sum()
        email_click_percentage = (email_clicks / total_students) * 100
        return email_click_percentage

email_click_percentage_apr = calculate_email_click_percentage(df,
    ↪ 'Email_click_Apr')
email_click_percentage_may = calculate_email_click_percentage(df,
    ↪ 'Email_click_May')

print(f'Percentage of students who clicked on the email in April:
    ↪ {email_click_percentage_apr:.2f}%')
print(f'Percentage of students who clicked on the email in May:
    ↪ {email_click_percentage_may:.2f}%')
```

Percentage of students who clicked on the email in April: 113.13%

Percentage of students who clicked on the email in May: 76.62%

Email clicks: The most email clicks were by first year students, while transfer students were second, but nowhere near the amount of clicks as the first years. I also discovered that April showed the highest average clicks per student. The story could be that first year students are more likely to engage with emails due to their unfamiliarity with college, while transfer students are less likely to engage with emails due to the fact they are more familiar with college processes.

Support for the Story:

```
[41]: df_first_year = df[df['Student Type'] == 'First Year Student']
df_transfer = df[df['Student Type'].str.contains('Transfer', na=False)]

email_click_percentage_first_year =
    ↪ calculate_email_click_percentage(df_first_year, 'Email_click_Apr')
email_click_percentage_transfer = calculate_email_click_percentage(df_transfer,
    ↪ 'Email_click_Apr')

print(f'Percentage of first year students who clicked on the email in April:
    ↪ {email_click_percentage_first_year:.2f}%')
print(f'Percentage of transfer students who clicked on the email in April:
    ↪ {email_click_percentage_transfer:.2f}%')
```

Percentage of first year students who clicked on the email in April: 113.55%

Percentage of transfer students who clicked on the email in April: 110.99%

Visuals:

```
[42]: plt.figure(figsize=(10,6))
sns.boxplot(x='Student Type', y='Email_click_Apr', data=df)
```

```

plt.title('April Email Clicks by Student Type')

plt.xticks(rotation=20)

plt.tight_layout()

plt.show()

plt.figure(figsize=(10,6))

sns.boxplot(x='Student Type', y='Email_click_May', data=df)


plt.title('May Email Clicks by Student Type')

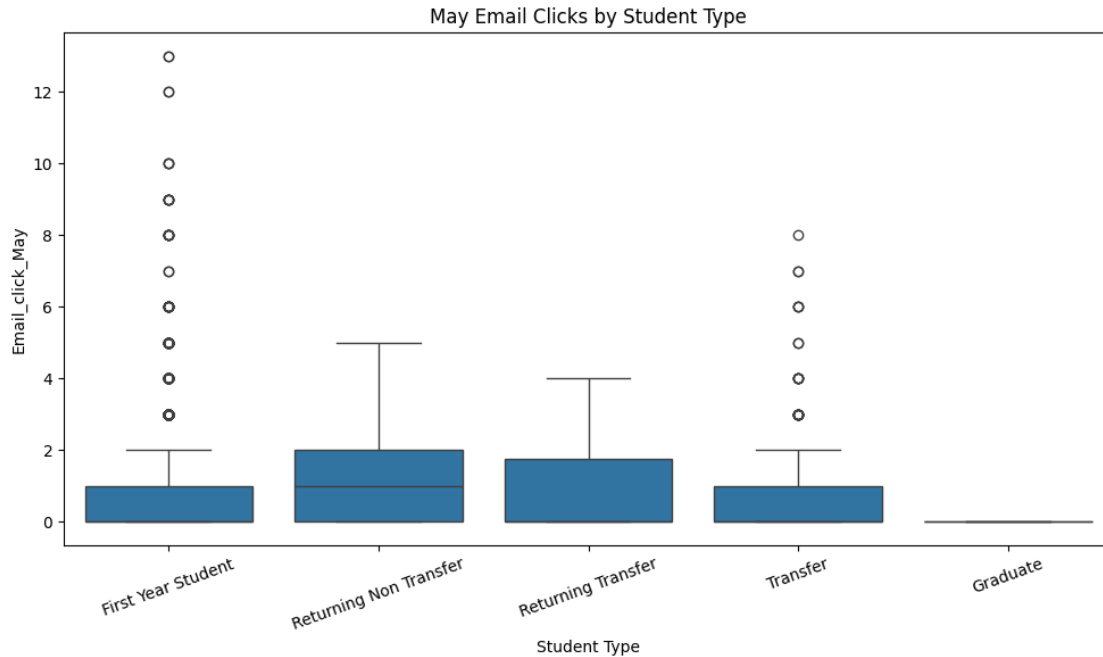
plt.xticks(rotation=20)

plt.tight_layout()

plt.show()

```





Step 2:

steps I would take to clean the dataset before analysis: 1. Handle missing values 2. Handle outliers 3. Normalize or standardize numerical data 4. categorical data 6. Feature selection

correlation between emails clicked and enrollment

```
[43]: df['Total_Clicks'] = df[['Email_click_Apr', 'Email_click_May',
    ↪ 'Email_click_Jun', 'Email_click_Jul', 'Email_click_Aug', 'Email_click_Sep',
    ↪ 'Email_click_Oct', 'Email_click_Nov', 'Email_click_Dec']].sum(axis=1)
df['Enrolled_Numeric'] = df['Enrolled'].map({'Yes': 1, 'No': 0})
correlation = df['Total_Clicks'].corr(df['Enrolled_Numeric'])
print(f"Correlation: {correlation}")
```

Correlation: 0.6308883699869055

This correlation shows the strength of the relationship between email engagement and enrollment. A positive correlation implies that higher email engagement is connected with a higher chance of enrollment.

```
[44]: df['Total_Engagement'] = df['Total_Clicks'] / df[['Email_click_Apr',
    ↪ 'Email_click_May', 'Email_click_Jun', 'Email_click_Jul', 'Email_click_Aug',
    ↪ 'Email_click_Sep', 'Email_click_Oct', 'Email_click_Nov', 'Email_click_Dec']].
    ↪ count(axis=1)

student_type_columns = [col for col in df.columns if col.startswith('Student_
    ↪ Type_')]
```

```

selected_features = student_type_columns + ['IsPastorKid', 'IsAthlete',
↪ 'Total_Clicks', 'Total_Engagement']
X = df[selected_features]
y = df['Enrolled_Numeric']
print("Available Types :", student_type_columns)
print("Selected features:", selected_features)

```

```

Available Types : []
Selected features: ['IsPastorKid', 'IsAthlete', 'Total_Clicks',
'Total_Engagement']

```

Which machine learning algorithms would I consider for this data set: - Logistic Regression: This algorithm is suitable for binary classification and it can handle both linear and non-linear relationships between features and the target. - K-nearest neighbors regression algorithm: Using this algorithm, we can find the nearest neighbors of a data point and make predictions based on the majority class or average value of those neighbors.

The model I chose is the K-nearest neighbor algorithm first

```

[45]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsRegressor

email_click_columns = [col for col in df.columns if col.
↪startswith('Email_click_')]
selected_features = ['IsPastorKid', 'IsAthlete'] + email_click_columns
X = df[selected_features].values
y = df['Enrolled'].map({'Yes': 1, 'No': 0}).values.astype(float)

X[:, 0] = (X[:, 0] == 'Yes').astype(int)
X[:, 1] = (X[:, 1] == 'Yes').astype(int)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↪random_state=42)

scaler = StandardScaler()
X_train_norm = scaler.fit_transform(X_train)
X_test_norm = scaler.transform(X_test)
k = 5
knn_regressor = KNeighborsRegressor(n_neighbors=k)
knn_regressor.fit(X_train_norm, y_train)
y_pred = knn_regressor.predict(X_test_norm)

mse = np.mean((y_pred - y_test) ** 2)
print(f"Mean Squared Error: {mse:.4f}")

```



```

r_squared = 1 - (np.sum((y_test - y_pred) ** 2) / np.sum((y_test - np.
↳mean(y_test)) ** 2))
print(f"R-squared: {r_squared:.4f}")

```

```

-----
ModuleNotFoundError                                Traceback (most recent call last)
Cell In[45], line 3
      1 import numpy as np
      2 import pandas as pd
----> 3 from sklearn.model_selection import train_test_split
      4 from sklearn.preprocessing import StandardScaler
      5 from sklearn.neighbors import KNeighborsRegressor

ModuleNotFoundError: No module named 'sklearn'

```

The K-nearest neighbor algorithm performed well with a Mean Squared Error of approximately 0.0929 and an R-squared value of approximately 0.6265. The model can accurately predict enrollment based on the given features.

Now I will test a Logistic Regression Algorithm

```

[ ]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)

from sklearn.linear_model import LogisticRegression

log_reg = LogisticRegression()

log_reg.fit(X_train, y_train)

y_pred = log_reg.predict(X_test)

from sklearn.metrics import accuracy_score

accuracy = accuracy_score(y_test, y_pred)

print(f"Accuracy: {accuracy:.4f}")

from sklearn.metrics import precision_score, recall_score, f1_score

precision = precision_score(y_test, y_pred)

recall = recall_score(y_test, y_pred)

f1 = f1_score(y_test, y_pred)

```

```
print(f"Precision: {precision:.4f}")
```

The Logistic Regression algorithm performed well with an accuracy of 0.9029. The precision was also high, showing that the model was able to correctly classify the enrollment status of students.

To improve the model, I would consider adding more features such as age, grade, and other demographic information.

Discuss potential biases in the dataset that could affect the model's predictions. - Selection bias: The dataset may not represent the diverse range of students in the college.

Ethical considerations when using this dataset include:

1. Bias in data collection: Ensure that the dataset is representative of the diverse range of students in the college.
2. Bias in data interpretation: Analyze the data carefully to ensure that it is not biased towards any specific group.
3. Bias in data analysis: Perform analysis on the dataset without making assumptions.

Privacy concerns when using this dataset include:

1. Data: Make sure that the dataset is securely stored and protected from unauthorized access.
2. Data misuse: don't share sensitive information without consent.
3. Data privacy: Make sure that the students' privacy is protected.